

Clustering concepts

Clustering of Data

Welcome to week 3

- By the end of this week you will know:
 - Distance metrics and their usage in machine learning.
 - Clustering concept and it's use for revealing patterns from unlabelled data

Distance Metrics

- Measuring **similarity or distances** between different data points is **fundamental to many machine learning algorithms**
 - unsupervised learning problems (i.e. K-means method in clustering)
 - supervised learning methods (i.e. K-Nearest-Neighbor)
- Distance measures are functions that **define a distance** $d(x_i, x_j)$ between any two data instances x_i and x_j **for measuring how similar** the instances are.


Distance Metrics...

- Distance measures satisfy the following three properties:
 - For any instance x_i , **distance with itself is zero**, $d(x_i, x_i) = 0$
 - For an instance pairs x_i and x_j , the **distance is non-negative and symmetric**,
$$d(x_i, x_j) \geq 0 \text{ and } d(x_i, x_j) = d(x_j, x_i)$$
 - Distance measure follows **triangular inequality** $d(x_i, x_k) \leq d(x_i, x_j) + d(x_j, x_k)$
- Distance measures satisfying above properties are also known as Distance Metrics

Distance Metrics...

- Example 1: Nearest Neighbour Classification

- Using distance to find the label of the new data point (the red triangle - Square or circle?)



- Example 2: Image retrieval

- Animal types in NUS Wide Animal dataset



- given a new image like the image of a cat, can we fetch all cat images from the dataset? (yes! with the help of distance measurements)

Distance Measurement Types

- Euclidean distance

- Straight-line distance between two points in Euclidean space
- For any two data instances, represented by **d-dimensional feature vectors** x_i, x_j their Euclidean distance is

$$d_{Euclidean}(x_i, x_j) = \left((x_{i,1} - x_{j,1})^2 + \dots + (x_{i,D} - x_{j,D})^2 \right)^{1/2}$$

- For example, consider these two vectors:

$$x_1 = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 1 \\ 0 \end{bmatrix} \text{ and } x_2 = \begin{bmatrix} 0 \\ 2 \\ 2 \\ 0 \\ 2 \end{bmatrix}$$

$$\begin{aligned} d(x_1, x_2) &= ((1 - 0)^2 + (1 - 2)^2 + (2 - 2)^2 + (1 - 0)^2 + (0 - 2)^2)^{\frac{1}{2}} \\ &= \sqrt{(1 + 1 + 0 + 1 + 4)} = \sqrt{7} = 2.65 \text{ (approx)} \end{aligned}$$

Distance Measurement Types

- Cosine distance

- For any two data instances, represented by d-dimensional feature vectors x_i, x_j

$$d_{\text{Cosine}}(x_i, x_j) = 1 - \frac{x_i^T x_j}{\|x_i\|_2 \cdot \|x_j\|_2}$$

- Lets see an example

$$d(x_1, x_2) = 1 - \frac{(1 * 0 + 1 * 2 + 2 * 2 + 1 * 0 + 0 * 2)}{\sqrt{1^2 + 1^2 + 2^2 + 1^2 + 0^2} * \sqrt{0^2 + 2^2 + 2^2 + 0^2 + 2^2}} \\ = 0.3453$$

$$x_1 = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 1 \\ 0 \end{bmatrix} \text{ and } x_2 = \begin{bmatrix} 0 \\ 2 \\ 2 \\ 0 \\ 2 \end{bmatrix}$$

Distance Measurement Types

- Cityblock/Manhattan distance

- For any two data instances, represented by d-dimensional feature vectors x_i, x_j , their Cityblock distance is computed as:

$$d_{\text{Cityblock}}(x_i, x_j) = |x_{i,1} - x_{j,1}| + \dots + |x_{i,D} - x_{j,D}|$$

- In most cases, this results similar to the Euclidean distance
- However, the effect of **a large difference in a single dimension is dampened** (since the distances are not squared)

$$d(x_1, x_2)_{CB} = |1 - 0| + |1 - 2| + |2 - 2| + |1 - 0| + |0 - 2| = 5$$

$$x_1 = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 1 \\ 0 \end{bmatrix}^T \quad \text{and} \quad x_2 = \begin{bmatrix} 0 \\ 2 \\ 2 \\ 0 \\ 2 \end{bmatrix}^T$$

Distance Measurement Types

- Minkowski distance

- Defines a distance between two points in a normed vector space.
- Euclidean and Cityblock distances are 2nd and 1st normed distance between x_i and x_j .
- Minkowski distance is a **generalization** of these distances defined for any p-norm

$$d(x, y) = \left(\sum_{i=0}^{n-1} |x_i - y_i|^p \right)^{\frac{1}{p}}$$

- when $p=1$, d is Manhattan distance
- when $p=2$, d is Euclidean distance

Distance Measurement Types

- Mahalanobis distance

- MD is the distance between two points in multivariate space, represented by d-dimensional feature vectors x_i, x_j

$$d_{Mahalanobis}(x_i, x_j) = (x_i - x_j)M^{-1}(x_i - x_j)^T$$

where, M is the covariance matrix of the data.

- Intuitively, the covariance matrix generalizes the notion of variance to multiple dimensions
- MD can be thought of **scaling each data dimension by its variance** and adjusting for their relationships
- When **data are independent**, i.e. $M=I$ (identity matrix), Mahalanobis distance becomes same as Euclidean distance.

Distance Measurement Types

- Jaccard distance

- Is a distance used to measure diversity of **any two sets**
- Consider any two instances x_i and x_j as **binary vectors indicating presence or absence of features**
- Jaccard distance between x_i and x_j is defined as $d_{Jaccard}(x_i, x_j) = 1 - \frac{|x_i \cap x_j|_1}{|x_i \cup x_j|_1}$
- Where \cap denotes logical 'AND' and \cup denotes logical 'OR' operators.
The $|x|_1$ is 1-norm
- Jaccard distance for $x_i = [1,0,1]$, and $x_j = [1,1,0]$

$$d_{Jaccard}(x_i, x_j) = 1 - \frac{1}{3} = \frac{2}{3}$$

Clustering and It's Applications


- Humans are encoded to see patterns in everything. (related to ML ?!)
- Did I just saw a huge puffy white duck in the sky?




- You are probably right! Our brain prefers patterns and we always look for patterns!
- It looks that our brains do clustering unconsciously

Clustering Algorithms

- How do we teach a computer to do this?
- Goal of clustering algorithms are:
 - Group objects of similar properties together
 - Discover interesting clusters and groups in the data
 - Find valid organisation of the data
- In other words, we can define two algorithmic goals:
 - Minimise intra-distance (distance between points in the same cluster)
 - Maximise inter-distance (distance between points from different clusters)



Clustering Algorithms...



- Now we can define a generic set-up based on our current understanding from clustering methods:
 - Step 1: define a distance metric between objects
 - Step 2: define an objective function that gets us to our clustering goal
 - Step 3: devise an algorithm to optimise the objective function

How Kmeans Works

- The most popular clustering algorithm; simple and fast
- was independently discovered in 60s and 70s by Steinhaus (1955), Lloyd (1957), Ball and Hall (1965) and McQueen (1967)
- Kmeans
 - stores k centroids
 - A point is considered to be in a particular cluster if it is closer to that cluster's centroid than any other centroid.
 - KMeans searches for the best centroids by alternating between two methods:
 - Assigning data points to clusters based on the current defined centroids (points which are the centre of a cluster).
 - Choosing centroids based on the current assignment of data points to clusters


N.B. Please see the video presented in “3.5 How Kmeans works” for better understanding

Evaluation of Clustering

- Generally there are two main categories of evaluation methods for clustering
- External assessment:
 - compare clustering performance against a known clustering (often called Ground truth or Gold standard)
- Internal assessment:
 - determine if clustering follows certain intrinsic assumptions (e.g. cluster-to-cluster distance or cluster size etc.),
 - Examples: Silhouette coefficient, Dunn index etc.

Evaluation of Clustering...

- The following figure illustrates a sample of ground truth (C) and the clustering partition found by a clustering algorithm (C')



Evaluation of Clustering...

- Rand Index

- Is a measure of the similarity between two data clusters
- Rand index is a function that measures the similarity of the two assignments C and C' , ignoring their permutations.
- Rand index is computed as
$$R = \frac{a + b}{\binom{n}{2}} = \frac{a + b}{a + b + c + d}$$
 - a = the number of pairs of data instances that are in the same cluster in both C, C'
 - b = the number of pairs of data instances that are in the different clusters in C and in different clusters in C'
 - c = the number of pairs of data instances that are in the same cluster in C but in different clusters in C' .
 - d = the number of pairs of data instances that are in the different clusters in C but in the same clusters in C'

Evaluation of Clustering...

- Purity

- Purity is a way of quality measurement in clustering methods
- As the name suggests, we would like to measure the purity for all clusters in terms of class labels of the data in each cluster i.e.,

$$Purity = \frac{\sum_{i=1..q} A_i}{\sum_{i=1..q} n_i}$$


where q is the number of clusters, A_i is the number of correctly assigned elements in i^{th} cluster and n_i is the number of elements in each cluster.

- This means the purity is measured by counting **the number of correctly assigned instances** and **dividing by the number of total instances**

Evaluation of Clustering...

- Purity Example

- Let's we have three clusters as shown in the figure.



- Each cluster is assigned to the class label which has the majority in the cluster i.e., first cluster labelled as Cross, second as Circle and last one as Plus.
- Based on the figure, 5 crosses, 4 circles, 3 pluses were correctly assigned. So, the purity of this clustering is:

$$Purity = \frac{1}{17} * (5 + 4 + 3) \approx 0.71$$

Evaluation of Clustering...

- Mutual Information

- Mutual information is a function that measures **the agreement of the two clustering assignments** C and C' in terms of how informative one is about the other, ignoring permutations.
- To put it simple, how informative is C about C'
- Let us assume that clustering partition C has K clusters and the partition C' has K' clusters

$$MI(C, C') = \sum_{i=1}^K \sum_{j=1}^{K'} P(i, j) \log \frac{P(i, j)}{P(i)P'(j)}$$

where $P(i)$ denotes the probability of randomly selected instance to belong to i^{th} cluster of the partition C, similarly, $P(i, j)$ denotes the probability of a randomly selected instance to belong in i^{th} cluster of the partition C and j^{th} cluster of the partition C'. $P'(j)$ is defined similarly as $P(i)$.

- So if our **C' clustering is highly informative based on C**, we can conclude that the C clustering assignment is doing good

Evaluation of Clustering...

- Silhouette Coefficient

- It is a measure of **how similar an object is to its own cluster** (cohesion) compared to other clusters (separation)
- This method does not require the ground truth cluster assignments
- The silhouette coefficient contrasts the average distance between the instances of the same cluster with the average distance between the instances of different clusters

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$

Where, $a(i)$ is the average distance of i^{th} instance with all other instances of the same cluster, and $b(i)$ is the lowest average dissimilarity of i^{th} instance with all other clusters.

- The final value of Silhouette ranges from -1 to $+1$. **High value of $s(i)$ indicates that the object is well matched to its own cluster and poorly matched to neighbouring clusters.**

Limitations of Kmeans

- Most important limitations of Simple Kmeans are:
 - Random initialisation means that you may get different clusters each time. As a solution, we can use Kmeans++ initialisation algorithm to initialise it better.
 - We have to supply the number of clusters beforehand. We can use Elbow method to choose K, but it may not be straightforward.
 - **It cannot find clusters of arbitrary shapes.**
 - It cannot detect noisy data points, i.e. they should not be taken into account for cluster analysis. K-median is less affected but cannot identify them.



Finding the Number of Clusters

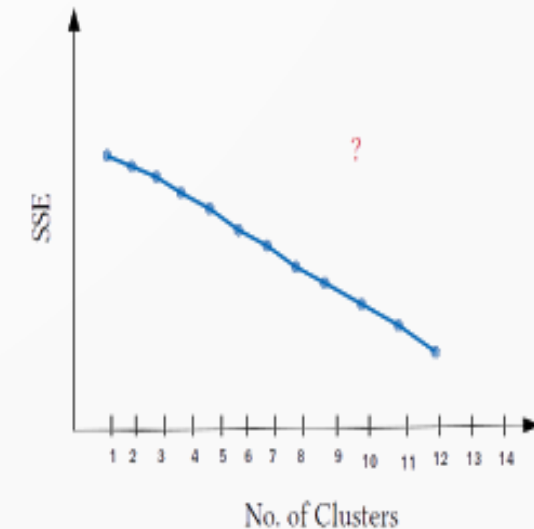
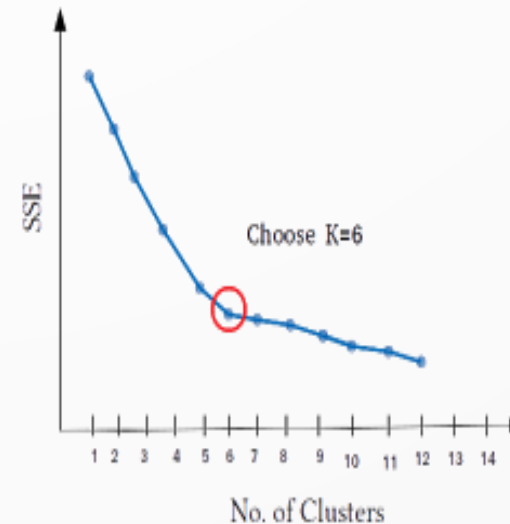
- Elbow Method

- The idea of elbow method is to run Kmeans clustering algorithm for a range of values of K, and for each value of K, compute the sum of squared error (SSE)

as:

$$SSE_k = \sum \sum z_{ik} \|x_i - \mu_k\|^2$$

- So, $\|x_i - \mu_k\|^2$ finds the distance of each point (x_i) to its corresponding centroid (μ_k) in the cluster
- z_{ik} is a binary variable
 - 1 when x_i is assigned to the cluster number k and
 - 0 when x_i is not related to cluster number k
- it looks 6 is the best cluster number



Kmeans with Kmeans++

- Kmeans++ is an algorithm for choosing the initial cluster's centre values or centroids for the Kmeans clustering algorithm
- K-means++ starts with allocating one cluster centre randomly and then searches for other centres given the first one
 - Choose one centroid μ_1 uniformly at random from dataset
 - Let $D(x)$ be the shortest distance from a data point to the closest centroid we have already chose
 - Choose a new centroid from the dataset with probability of $\frac{D^2(x_i)}{\sum_i D^2(x_i)}$
 - Now repeat previous step until we have initialised K centroids

Kmeans with Kmeans++

- Guarantee of Kmeans++

- In Kmeans algorithm with random initialisation of centroids, the objective function monotonically decreases with each iteration of the algorithm
- Let us say, for the best solution, the objective function takes value $j_{optimum}$
- Let us say, when using Kmeans the objective function converges to $\hat{j}_{optimum}$
- For Kmeans with random initialisation there is no theoretical bound on
$$\frac{\hat{j}_{converged}}{\hat{j}_{optimum}}$$
- In contrast, Kmeans++ initialisation has the following theoretical guarantee on convergence

$$\frac{E(\hat{j}_{converged})}{\hat{j}_{optimum}} \leq 8(\log K + 2)$$

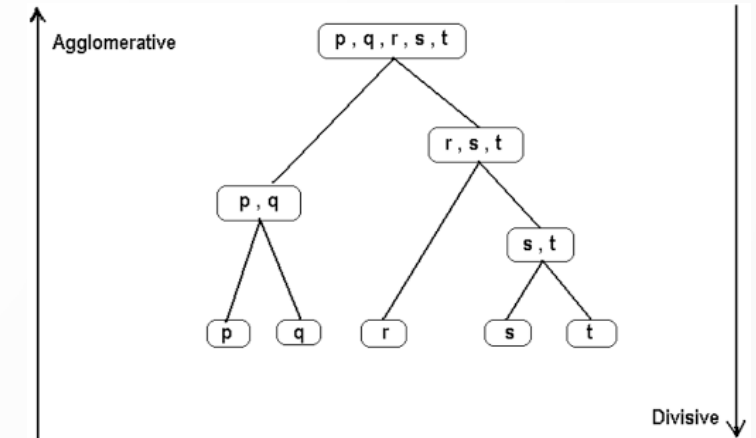
Other Clustering Algorithms...

- Hierarchical Clustering:
 - clusters that have a predetermined ordering
 - Two types -
 - **Agglomerative Clustering (Bottom-up)**
 - each observation starts in its own cluster, and pairs of clusters are merged as one moves up the hierarchy
 - **Divisive Clustering (Top-down)**
 - all observations start in one cluster, and splits are performed recursively as one moves down the hierarchy

Other Clustering Algorithms...

- Agglomerative Clustering:

- At the bottom of the tree, at the starting point each of the characters p,q,r,s,t are assigned into a single separate cluster
- As we go up to the higher levels, the closest characters are formed another cluster. i.e. s,t and p,q.
- At the next level we can notice r,s,t are making a cluster
- And finally at the top of the tree, all the characters are in one single cluster p,q,r,s,t
- How to find the closest cluster pairs? i.e. how to find the distance between two sub-clusters in the middle of the tree?



Other Clustering Algorithms...

- Agglomerative Clustering: four ways to find distance -
 - Single-link: It is a distance between closest points
 - Complete-link: Distance between the furthest points
 - Centroid: Distance between the Centroids
 - Average-link: Average distance between pairs of elements from across cluster pairs

Other Clustering Algorithms...

- Divisive Clustering

- Same as the Agglomerative Clustering in this type of clustering, initially all data instances are put in the same cluster
- For splitting, we can use any clustering algorithm that produces at least two clusters (e.g. Kmeans) to find 2 clusters
- The process is continued until each data instance is separate

Thank You.