

Model Assessment & Selection

Week 6

Model Selection

Partitioning data for train and test


Mean square error

Finding the best hyper-parameters

Effects of Imbalanced class

Partitioning data for training and testing

- Reliable estimate of model performance (accuracy)
 - Training/test split of the data
 - Single
 - Multiple
 - Larger test set
 - Larger training set

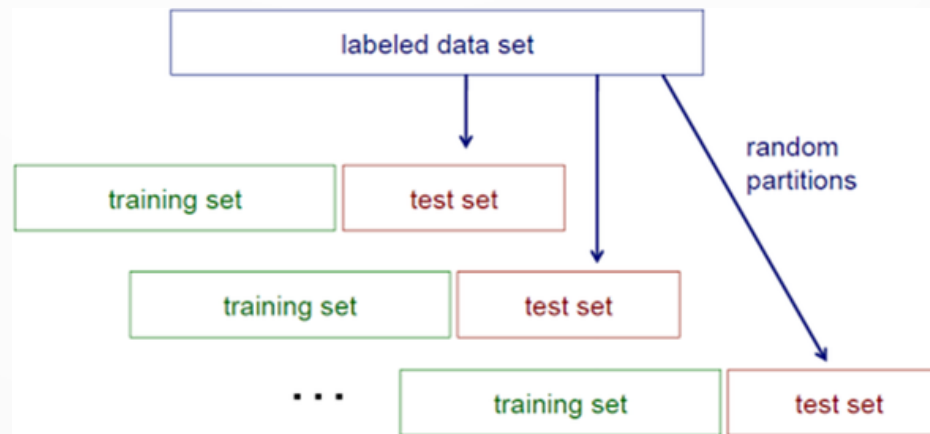


Partitioning data for training and testing

- Methods for splitting data:
 - Random sub-sampling
 - Stratified sampling
 - Cross validation

Partitioning data: Random sub-sampling

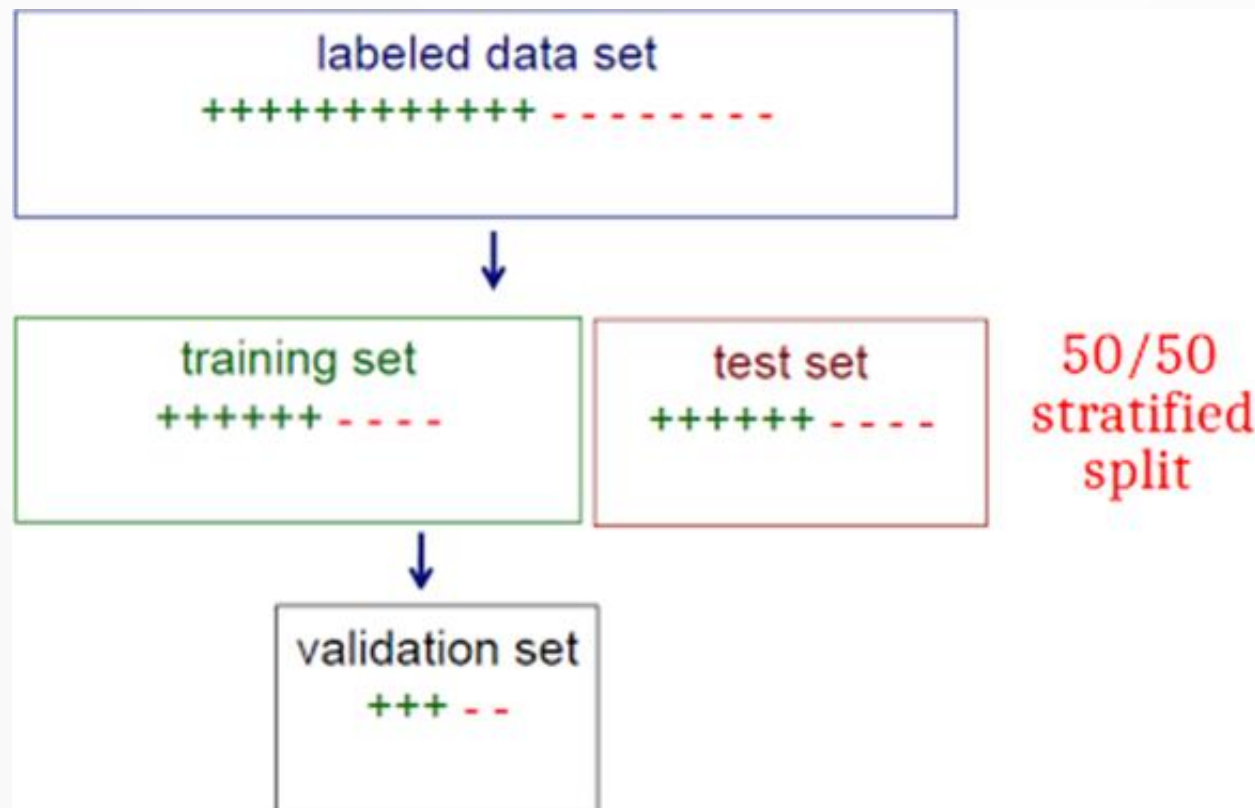
- A more reliable estimate of model performance can be obtained by **random sub-sampling**.
- Random sub-sampling
 - Partitions the data into random training and test sets in a specified ratio.



- For **multiple splits / trials**
 - Average the accuracies to get an averaged estimate.

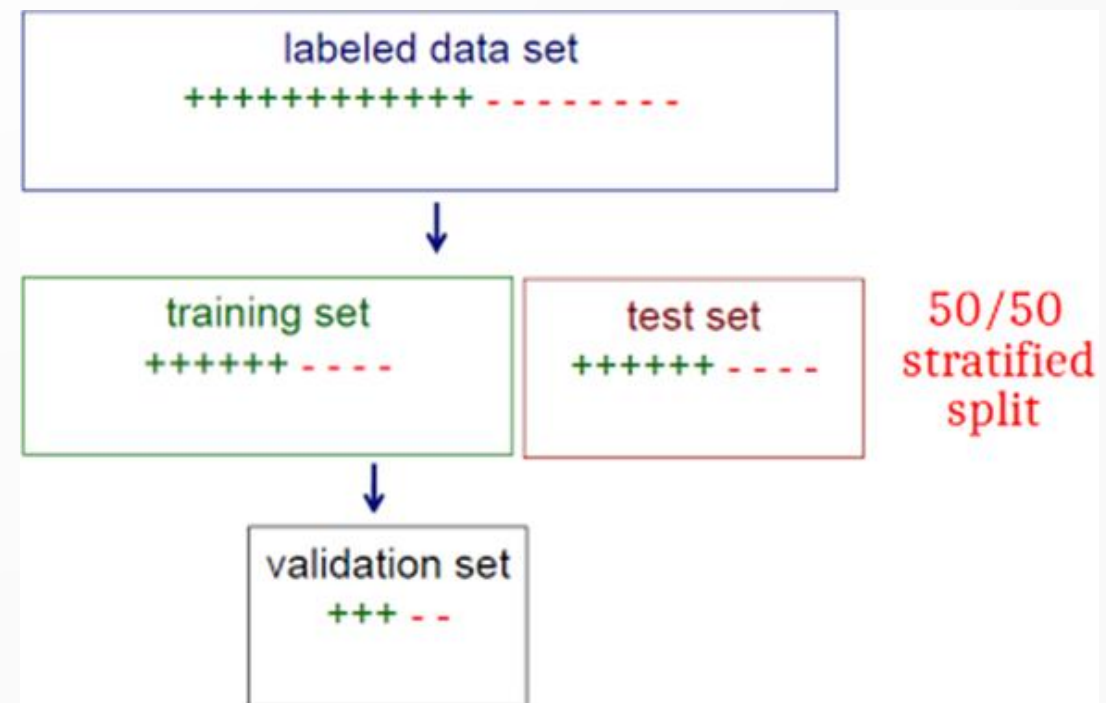
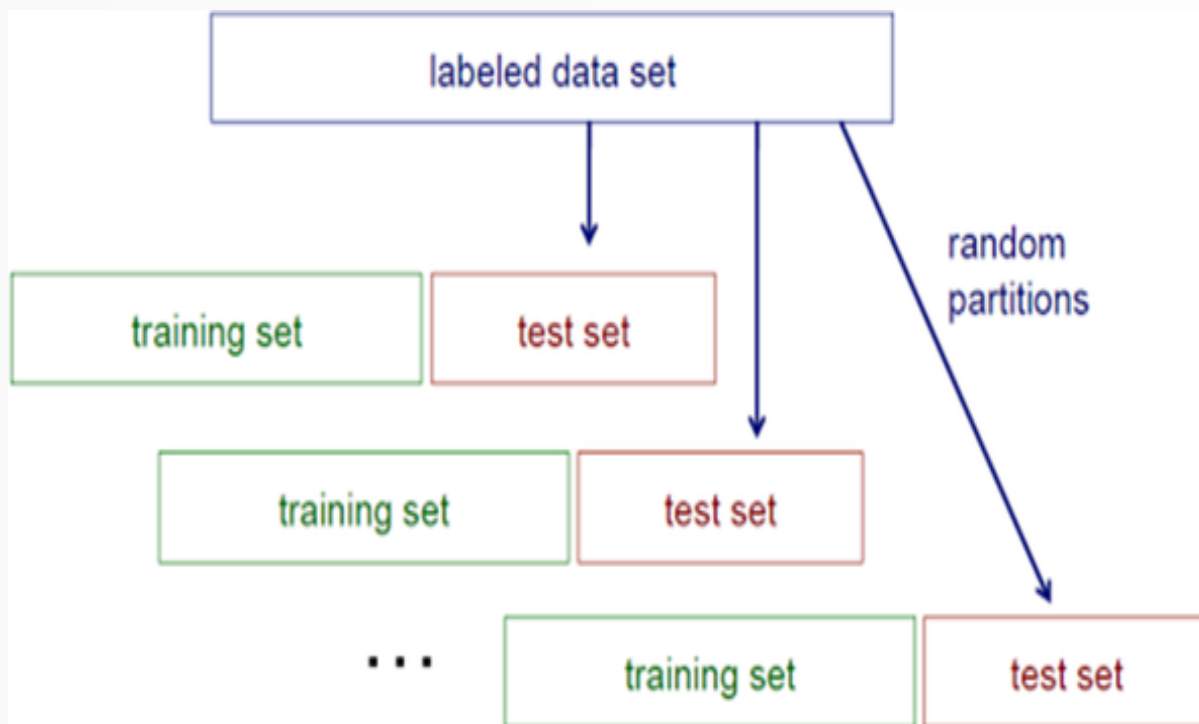
Partitioning data: Stratified Sampling

- Stratified sampling is a probability sampling technique
 - Divide the entire data into different subgroups or strata.
 - Randomly selects the samples proportionally from the different strata.





Partitioning data: Stratified Sampling...

- Difference with random sampling:
 - When randomly selecting training (or validation) sets, **class proportions may differ** between training and test splits.



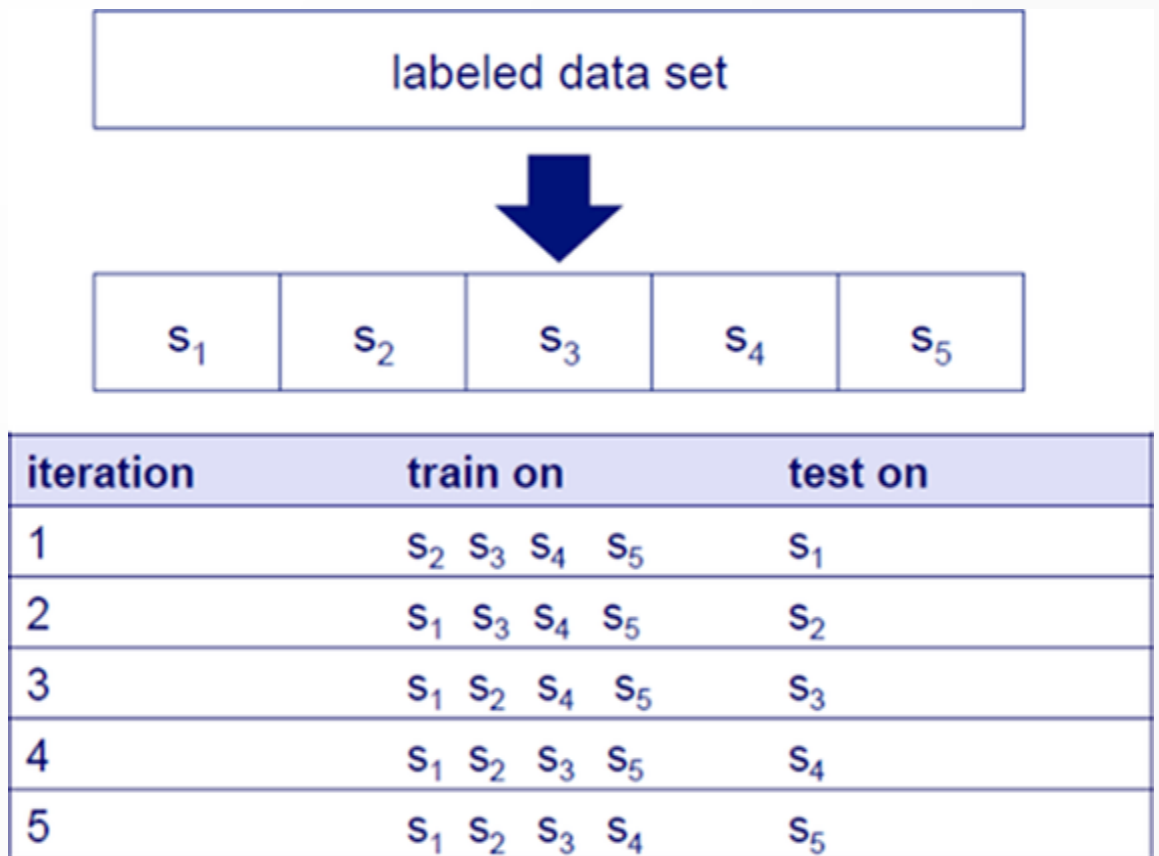
Partitioning data: Stratified Sampling...

- For **multi-class** problem:



Partitioning data: Cross-validation

- Cross-validation.
 - Evaluate models by **partitioning the original sample** into a **training set** to train the model, and a **test set** to evaluate it.
 - The main idea is **to partition training data into k equal sized sub-samples (k-fold cross-validation)**.
 - Iteratively leave **one sub-sample out** for evaluating the model and train the model using **the rest of the sub-samples**.



Partitioning data: Cross-validation...

- In summary:
 - Accuracy can be averaged over multiple runs. (sounds like multiple splitting!)
 - In the special case, **when k is equal to the number of instances n ,**
 - **Leave-one-out cross-validation** scheme.
 - Cross-validation makes efficient use of the available data for testing.

Iteration	Train on	Test on	Correct
1	S_2, S_3, S_4, S_5	S_1	110/200
2	S_1, S_3, S_4, S_5	S_2	170/200
3	S_1, S_2, S_4, S_5	S_3	160/200
4	S_1, S_2, S_3, S_5	S_4	130/200
5	S_1, S_2, S_3, S_4	S_5	160/200

Accuracy = 730/1000 = 73%

Mean Square Error in Linear Regression

- Mean Square Error (MSE) is a metric used to evaluate linear models


$$MSE = \frac{1}{n} \sum_{i=1}^n (\hat{Y}_i - Y_i)^2$$

- The goal of any model should be to reduce the MSE
 - Why?
 - A **smaller MSE** implies that there is relatively **little difference** between the estimated and observed outputs.
 - A well-fitted model should have a relatively low MSE value.
 - The ideal form has an **MSE of zero**
 - **No difference** between the estimated and observed parameters.

MSE in Linear Regression...

The effects of test size on MSE

- There are two cases in which you have to be extremely careful when using MSE to compare your results with the true outputs.
- Case-1: Anscombe's quartet comprises four datasets that have **nearly identical** simple descriptive statistics, yet appear very different when graphed.
- MSE is not a wise measurement in all these cases except in set 1 (**as the MSE values will be the same in all 4 cases**)



MSE in Linear Regression: The effects of test size on MSE...

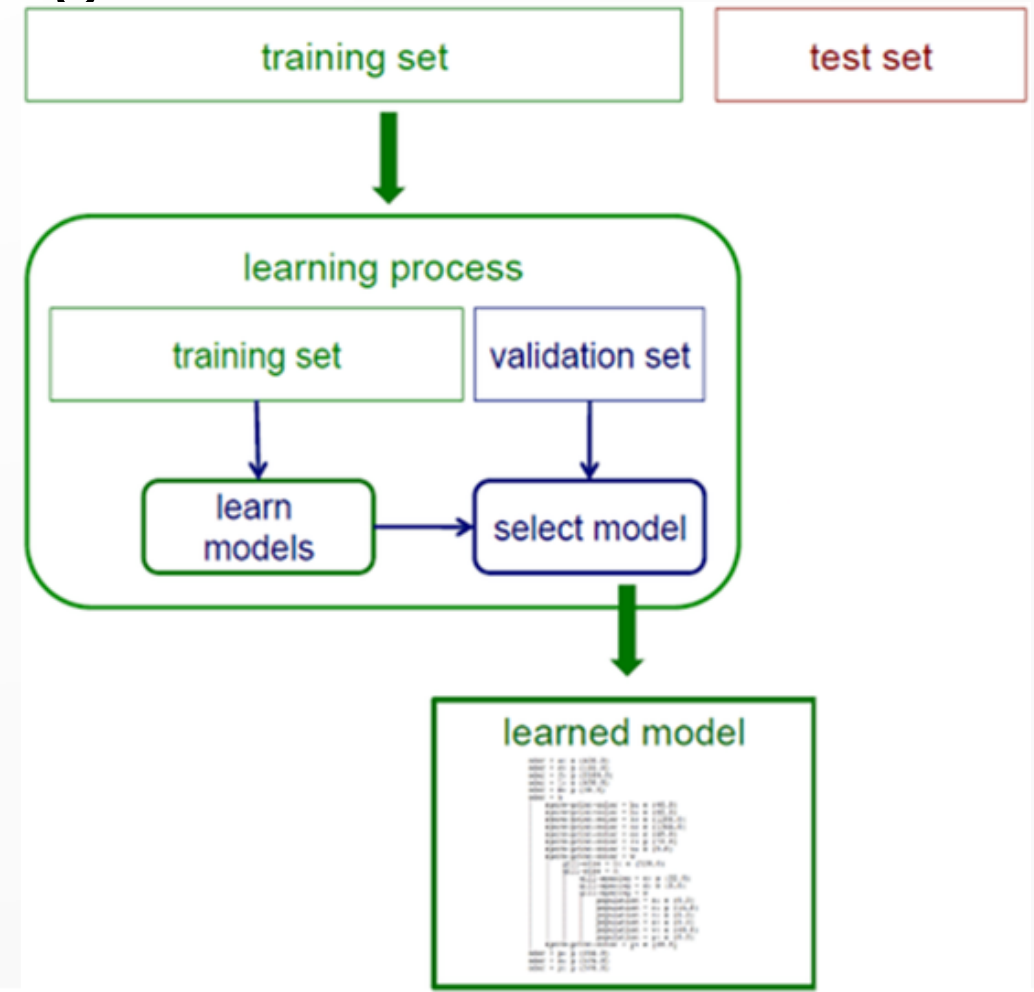
- Case-2: The **number of test data points** while calculating MSE of your model is important.
 - What if you have two models and you want to compare their performance based on MSE.
 - One of these models has 100 test points and the other one has been evaluated on 100000 test points.
 - Is it fair to compare such models?
 - Evaluate different models with the same number of test data points.

Finding the best hyper-parameters

- What is a hyperparameter?
 - a parameter whose **value is set** before the learning process begins.
 - the value of a hyperparameter in a model **cannot be estimated from data**.
 - often used in processes to **help estimate model parameters**.
 - often set by **using heuristics**
 - often **tuned** for a given predictive modeling problem

Finding the best hyper-parameters...

- To search for the best hyperparameters, we need to partition training data into separate training and validation sets
- We already know about training and test data.
- But what is **validation set**?



Finding the best hyper-parameters...

- A validation set:
 - a sample of data used to provide an unbiased evaluation of a model fit on the training dataset
 - fine-tune the model hyperparameters. How?
 - evaluate the performance of the model for different combinations of hyperparameter values (e.g. by means of a grid search process) and keep the best trained model or hyperparameters.
 - Why validation set?
 - Keeping test set unbiased to model training or hyperparameter selection
 - Since test set is used for comparing different models

Finding the best hyper-parameters...

- But, how can we exactly find the best hyper-parameter?
 - First, we need to decide a possible range for hyperparameter, i.e., a bounded interval such as $[0,1]$
 - We then define a search grid within the specified range. i.e., we would like to select these values $\{0, 10^{-3}, 10^{-2}, 10^{-1}, 1\}$ for hyperparameter in order to evaluate the model with them.
 - Next, we train a model using each hyperparameter value from the search grid and assess its performance on a validation set (taken out from the training set).
 - Finally, We compute the performance on the validation set for each hyper-parameter value and select the one with the best performance.

Finding the best hyper-parameters...

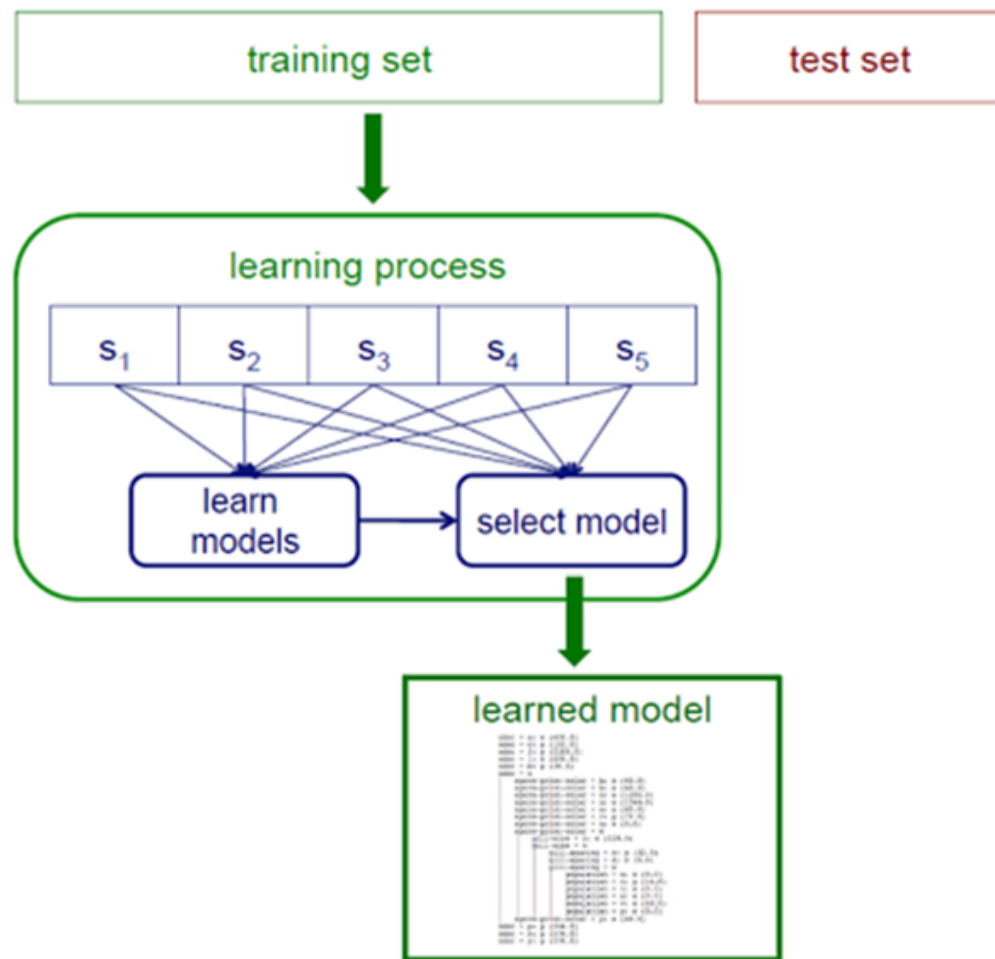
- Note: In the above example we just considered 4 cases to evaluate as a hyperparameter in the $[0,1]$ interval.
- In this continuous space, we may lose many other good options by restricting the search only to 4 values.
- It is obvious in this case we will be more accurate in finding the value of the hyperparameter.
- But this Grid-searching can be extremely computationally expensive

Finding the best hyper-parameters: Internal cross-validation

- All the techniques that we previously **discussed for model assessment are applicable for training/validation** set splitting:
 - Random subsampling
 - Stratified subsampling
 - Cross-validation
- Remember, this **step is internal to the learning process** and different from model assessment on the test data.

Finding the best hyper-parameters: Internal cross-validation...

- Let us examine how an internal cross-validation works.



Finding the best hyper-parameters: Internal cross-validation...

- Remember that we can select **the best hyperparameter set** by searching/or optimizing over all possible values.
- Let us show you 3 possible ways to navigate the hyperparameter space:
 - **Grid-search** (not so efficient). This is what we are using and explaining!
 - **Random search** (efficient in certain scenarios) [Bergstra et al. (JMLR 2012)]
 - **Bayesian optimization** (efficient in general) [Snoek et al. (2012)]


Finding the best hyper-parameters:

Effect of imbalanced class

- What is imbalanced data set?
 - Ratio of class distribution is skewed toward some class/es.
 - the total number of a class of data (i.e. positive) is far less than the total number of another class of data (i.e. negative).
- This problem is very common in practice
 - fraud detection
 - anomaly detection
 - medical diagnosis, etc.
- Most machine learning algorithms work best when the number of instances of each classes are roughly equal.

Finding the best hyper-parameters: Effect of imbalanced class

- Example: When developing a breast cancer diagnosis model, imbalanced class problem is encountered because the risk of a female being diagnosed with breast cancer (by their 85th birthday) is 1 in 8.
- This means that a representative training set will have 7 times more instances in negative class than the positive class (see figure below)



Finding the best hyper-parameters:

Effect of imbalanced class

- So what are the possible solutions to overcome this problem?
 - Perform some actions **on the data** itself.
 - **Improve the algorithm** to be able to handle such phenomenon.
- At the data level: (Re-Sampling):
 - over-sampling the data from minority class
 - under-sampling the data from majority class.
- At the algorithmic level:
 - adjusting the costs
 - adjusting the decision threshold.

Finding the best hyper-parameters:

Issues of imbalanced classes

- Now, let us have a close look on possible issues of imbalanced classes.
- Problem-1: Since the test data contains **only few samples from the minority class**, even a dumb classifier that always classifies an instance to the majority class will get very high accuracy!
 - This problem is dealt by using **other evaluation metrics in place of accuracy**.
- Problem-2: When doing **random subsampling**, it is possible that class proportion is not maintained in individual partition. In fact, we may not sample even one instance from the minority class.
 - This problem can be solved using **Stratified Sampling**.

Finding the best hyper-parameters:

Issues of imbalanced classes

- But always remember:
 - Any **pre-processing over entire data set** (e.g. feature selection, or feature extraction) must not use the information that you are trying to predict (e.g. labels).
 - In the training process, you must not use any information that is not available during the training process.
 - Example: I was building a cancer prognosis model to predict whether a patient will survive 1 year from diagnosis or not?
 - I used the cause of death field as one of the features.
 - Clearly, this information is not available at the prediction time.
 - If you **modify your model again and again** by looking at how it performs on a specified test set, then you may be overfitting on the test set.

Thank You.